

Dimension Reduction

Data Mining for Business Analytics in Python

Shmueli, Bruce, Gedeck & Patel

Dimension Reduction

- The dimension of a dataset, which is the number of variables, must be reduced for the data mining algorithms to operate efficiently. This process is part of the pilot/prototype phase of data mining and is done before deploying a model.
- Several dimension reduction approaches:
 - 1) Incorporating domain knowledge to remove or combine categories
 - 2) Using data summaries to detect information overlap between variables (and remove or combine redundant variables or categories)
 - 3) Using data conversion techniques such as converting categorical variables into numerical variables
 - 4) Employing automated reduction techniques, such as Principal Components Analysis (PCA), where a new set of variables (which are weighted averages of the original variables) is created. These new variables are uncorrelated and a small subset of them usually contains most of their combined information (hence, we can reduce dimension by using only a subset of the new variables).
- Introducing data mining methods such as “regression models and classification” and “regression trees”, which can be used for removing redundant variables and for combining “similar” categories of categorical variables.
- In Artificial Intelligence literature, dimension reduction is often referred to as *factor selection* or *feature extraction*.

Exploring the data

Statistical summary of data: common metrics

- Average
- Median
- Minimum
- Maximum
- Standard deviation
- Counts & percentages

Exploring the data

```
bostonHousing_df = pd.read_csv('BostonHousing.csv')  
bostonHousing_df = bostonHousing_df.rename(columns={'CAT.  
MEDV': 'CAT_MEDV'})
```

```
# Compute mean, standard dev., min, max, median, length,  
and missing values for all variables
```

```
pd.DataFrame({'mean': bostonHousing_df.mean(),  
'sd': bostonHousing_df.std(),  
'min': bostonHousing_df.min(),  
'max': bostonHousing_df.max(),  
'median': bostonHousing_df.median(),  
'length': len(bostonHousing_df),  
'miss.val': bostonHousing_df.isnull().sum(),  
})
```

Summary Statistics for Boston Housing Data

	mean	sd	min	max	median	length	miss.val
CRIM	3.61352356	8.6015451	0.00632	88.9762	0.25651	506	0
ZN	11.36363636	23.3224530	0.00000	100.0000	0.00000	506	0
INDUS	11.13677866	6.8603529	0.46000	27.7400	9.69000	506	0
CHAS	0.06916996	0.2539940	0.00000	1.0000	0.00000	506	0
NOX	0.55469506	0.1158777	0.38500	0.8710	0.53800	506	0
RM	6.28463439	0.7026171	3.56100	8.7800	6.20850	506	0
AGE	68.57490119	28.1488614	2.90000	100.0000	77.50000	506	0
DIS	3.79504269	2.1057101	1.12960	12.1265	3.20745	506	0
RAD	9.54940711	8.7072594	1.00000	24.0000	5.00000	506	0
TAX	408.23715415	168.5371161	187.00000	711.0000	330.00000	506	0
PTRATIO	18.45553360	2.1649455	12.60000	22.0000	19.05000	506	0
LSTAT	12.65306324	7.1410615	1.73000	37.9700	11.36000	506	0
MEDV	22.53280632	9.1971041	5.00000	50.0000	21.20000	506	0
CAT.MEDV	0.16600791	0.3724560	0.00000	1.0000	0.00000	506	0

Correlation Matrix for Boston Housing Data

portion of the correlation table shown, using `corr` from `pandas`

```
> bostonHousing_df.corr().round(2)
```

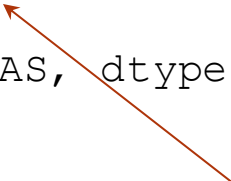
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CRIM	1.00	-0.20	0.41	-0.06	0.42	-0.22	0.35	-0.38
ZN	-0.20	1.00	-0.53	-0.04	-0.52	0.31	-0.57	0.66
INDUS	0.41	-0.53	1.00	0.06	0.76	-0.39	0.64	-0.71
CHAS	-0.06	-0.04	0.06	1.00	0.09	0.09	0.09	-0.10
NOX	0.42	-0.52	0.76	0.09	1.00	-0.30	0.73	-0.77
RM	-0.22	0.31	-0.39	0.09	-0.30	1.00	-0.24	0.21
AGE	0.35	-0.57	0.64	0.09	0.73	-0.24	1.00	-0.75
DIS	-0.38	0.66	-0.71	-0.10	-0.77	0.21	-0.75	1.00

Using `value_counts()`, in pandas, to tabulate counts

```
> bostonHousing_df.CHAS.value_counts()
```

```
0    1  
471 35
```

```
Name: CHAS, dtype: int64
```



35 neighborhoods have a CHAS
value of "1," i.e. they border the
Charles River

Using `groupby()` to tabulate counts using multiple variables

```
# Create bins of size 1 for variable using the method pd.cut. Default
# creates a categorical variable, e.g. (6,7]. labels=False determines
# integers instead, e.g. 6.
```

```
bostonHousing_df['RM_bin'] = pd.cut(bostonHousing_df.RM, range(0, 10), labels=False)
```

```
# Compute average of MEDV by (binned) RM and CHAS. First group the data frame
# using the groupby method, then restrict the analysis to MEDV and determine the
# mean for each group.
```

```
bostonHousing_df.groupby(['RM_bin', 'CHAS'])['MEDV'].mean()
```

RM_bin	CHAS	
3	0	25.300000
4	0	15.407143
5	0	17.200000
	1	22.218182
6	0	21.769170
	1	25.918750
7	0	35.964444
	1	44.066667
8	0	45.700000
	1	35.950000

← In neighborhoods where houses averaged 3 rooms and did not border the Charles, median value was 25.3 (\$000)

RM = Average number of rooms per dwelling

- `pd.cut(bostonHousing_df.RM, range(0, 10), labels=False)`: This part of the code applies the `pd.cut` function to the 'RM' column. The `pd.cut` function is used to create discrete bins for the values in the specified column.
- `bostonHousing_df.RM`: This is the Series containing the values to be binned.
- `range(0, 10)`: This specifies the bin edges. It creates bins from 0 (inclusive) to 10 (exclusive). So, you'll have bins like [0, 1), [1, 2), [2, 3), and so on.
- `labels=False`: This parameter specifies that you want the output to be numeric bin labels rather than categorical labels.

Use `pivot_table()` in pandas for pivot tables

use `pivot_table()` to reshape data and generate pivot table

```
pd.pivot_table(bostonHousing_df, values='MEDV', index=['RM_bin'],
               columns=['CHAS'], aggfunc=np.mean, margins=True)
```

CHAS	0	1	All
RM_bin			
3	25.300000	NaN	25.300000
4	15.407143	NaN	15.407143
5	17.200000	22.218182	17.551592
6	21.769170	25.918750	22.015985
7	35.964444	44.066667	36.917647
8	45.700000	35.950000	44.200000
All	22.093843	28.440000	22.532806

neighborhoods not bordering Charles River, with 8 rooms, have MEDV = 45.7

- `aggfunc=np.mean`: This specifies that the aggregation function to be used is `np.mean`, which calculates the mean (average) of values in each group of the pivot table. We do not need to insert `aggfunc=np.mean` as “avg” is the default parameter for `aggfunc`.
- `margins=True`: This includes the margins (totals) in the pivot table, giving you the mean values for 'MEDV' across all 'RM_bin' and 'CHAS' groups.
- **margins** *bool, default False* If `margins=True`, special All columns and rows will be added with partial group aggregates across the categories on the rows and columns.


Reducing Categories

- A single categorical variable with m categories is typically transformed into m or $m-1$ dummy variables
- Each dummy variable takes the values 0 or 1
 - 0 = “no”
 - 1 = “yes”
- Problem: Can end up with too many variables
- Solution: Reduce by combining categories that are close to each other
- Use pivot tables to assess outcome variable sensitivity to the dummies
- Exception: Naïve Bayes can handle categorical variables without transforming them into dummies

Combining Categories

Stacked bar chart:

- Many zoning categories are the same or similar with respect to CATMEDV



ZN = Percentage of residential land zoned for lots over 25,000 ft²

CAT.MEDV = Is median value of owner-occupied homes in tract above \$30,000 (CAT.MEDV = 1) or not (CAT.MEDV = 0)

Example: For a lot with 33 percent of residence, 50% of homes belong to category 0 and 50% of home belong to category 1

CODE FOR STACKED BAR CHART

```
# use method crosstab to create a cross-tabulation of two  
variables
```

```
tbl = pd.crosstab(bostonHousing_df.CAT_MEDV, bostonHousing_df.ZN)
```

```
# convert numbers to ratios
```

```
propTbl = tbl / tbl.sum()
```

```
propTbl.round(2)
```

```
# plot the ratios in a stacked bar chart
```

```
ax = propTbl.transpose().plot(kind='bar', stacked=True)
```

```
ax.set_yticklabels(['{:, .0%}'.format(x) for x in
```

```
ax.get_yticks()])
```

```
plt.title('Distribution of CAT.MEDV by ZN')
```

```
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
```

```
plt.show()
```

- The format string '{:, .0%}' is used to format a numeric value as a percentage with specific formatting:
- : Marks the beginning of the format specification.
- ,: Specifies that a comma should be used as a thousand separator for large numbers. For example, the number 1000 would be formatted as "1,000."
- .0: Specifies that there should be no decimal places in the formatted percentage.
- %: Specifies that the value should be presented as a percentage, and the % symbol will be appended to the end. So, it will display percentages like "50%" or "75%" with a comma for thousands if needed.
- The legend appears to the right of the plot because `bbox_to_anchor=(1, 0.5)` places the legend's anchor point at the right edge of the axes. The `loc='center left'` specifies that the center of the left side of the legend box aligns with that anchor point, thus moving the legend outside the plot area on the right.

Principal Components Analysis

Goal: Reduce a set of **numerical** variables.

The idea: Remove the overlap of information between these variable. ["Information" is measured by the sum of the variances of the variables.]

Final product: A smaller number of numerical variables that contain most of the information

Principal Components Analysis

How does PCA do this?

- Creates new variables that are linear combinations of the original variables (i.e., they are weighted averages of the original variables).
- These linear combinations are uncorrelated (no information overlap), and only a few of them contain most of the original information.
- The new variables are called *principal components*.

Example – Breakfast Cereals (excerpt)

name	mfr	type	calories	protein	...	rating
100%_Bran	N	C	70	4	...	68
100%_Natural_Bran	Q	C	120	3	...	34
All-Bran	K	C	70	4	...	59
All-Bran_with_Extra_Fiber	K	C	50	4	...	94
Almond_Delight	R	C	110	2	...	34
Apple_Cinnamon_Cheerios	G	C	110	2	...	30
Apple_Jacks	K	C	110	2	...	33
Basic_4	G	C	130	3	...	37
Bran_Chex	R	C	90	2	...	49
Bran_Flakes	P	C	90	3	...	53
Cap'n'Crunch	Q	C	120	1	...	18
Cheerios	G	C	110	6	...	51
Cinnamon_Toast_Crunch	G	C	120	1	...	20

Description of Variables

name: name of cereal

mfr: manufacturer

type: cold or hot

calories: calories per serving

protein: grams

fat: grams

sodium: mg

fiber: grams

carbo: grams complex
carbohydrates

sugars: grams

potass: mg

vitamins: % FDA rec

shelf: display shelf

weight: oz. 1 serving

cups: # in one serving

rating: consumer reports

Consider calories & ratings covariance matrix

$$\text{Covariance Matrix Formula: } \begin{bmatrix} \text{Var}(X_1) & \dots & \text{Cov}(X_n, X_1) \\ \vdots & \ddots & \vdots \\ \text{Cov}(X_n, X_1) & \dots & \text{Var}(X_n) \end{bmatrix}$$

```
cereal_df = pd.read_csv('Cereals.csv')  
cereal_df[['calories', 'rating']].cov()
```

```
cereal_df.calories.var()  
cereal_df.rating.var()
```


	calories	ratings
calories	379.63	-189.68
ratings	-189.68	197.32

- Total variance (=“information”) is sum of individual variances: $379.63 + 197.32 = 577$
- Calories accounts for $379.63/577 = 66\%$
- If we want to make decision with just calories, we lose 34% of the variation
- The main purpose of a covariance matrix is to show how pairs of variables in a dataset covary or move together. A positive covariance indicates that when one variable increases, the other tends to increase as well, while a negative covariance indicates that one variable tends to decrease when the other increases.
- The magnitude of the covariances can indicate the degree of dependence or independence between variables. If the covariances are close to zero, the variables are relatively independent. Larger positive or negative covariances indicate stronger relationships.
- While the covariance provides information about the direction and strength of the linear relationship between variables, it does not offer a standardized measure like correlation. Correlation, which is derived from the covariance matrix, provides a normalized measure that ranges between -1 and 1, making it easier to interpret.

Using linear combinations to redistribute the variability in a more polarized way

Z_1 and Z_2 are two linear combinations.

- Z_1 has the highest variation (spread of values)
- Z_2 has the lowest variation



PCA output for these two variables

```
from sklearn.decomposition import PCA
pcs = PCA(n_components=2)
pcs.fit(cereals_df[['calories', 'rating']])
```

Weights to project original data onto Z_1 & Z_2 , e.g. (-0.847, 0.532) are weights for Z_1

```
pcs.explained_variance_ratio_
pcs_df = pd.DataFrame(pcs.components_, columns=['PC1', 'PC2'])
```

	PC1	PC2
cereals.df.calories	-0.8470535	0.5315077
cereals.df.rating	0.5315077	0.8470535

```
pcsSummary = pd.DataFrame({'Standard deviation': np.sqrt(pcs.explained_variance_),
                           'Proportion of variance': pcs.explained_variance_ratio_,
                           'Cumulative proportion': np.cumsum(pcs.explained_variance_ratio_)})
pcsSummary = pcsSummary.transpose()
```

Importance of components:

	PC1	PC2
Standard deviation	22.3165	8.8844
Proportion of Variance	0.8632	0.1368
Cumulative Proportion	0.8632	1.0000

86% of the total variance is accounted for by component 1

Principal Component Scores for the First Five Records

$\text{Score_PC1} = \text{Weight_PC1} * (\text{Value_PC1} - \text{Mean of Column/Variable in PC1})$

$\text{Score_PC2} = \text{Weight_PC2} * (\text{Value_PC2} - \text{Mean of Column/Variable in PC2})$

$\text{SCORE} = \text{Score_PC1} + \text{Score_PC2}$

```
pd.DataFrame(pcs.transform(cereal_df[['calories', 'rating'])),  
              columns=[PC1', 'PC2'])
```

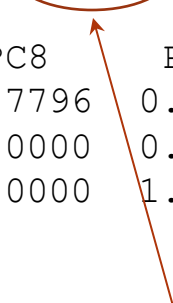
	PC1	PC2
[1,]	44.921528	2.1971833
[2,]	-15.725265	-0.3824165
[3,]	40.149935	-5.4072123
[4,]	75.310772	12.9991256
[5,]	-7.041508	-5.3576857

PCA for the 13 Numerical Variables in the Cereals Data

```
pcs = PCA()
pcs.fit(cereals_df.iloc[:, 3:].dropna())
pcsSummary_df = pd.DataFrame({'Standard deviation':
    np.sqrt(pcs.explained_variance_), 'Proportion of variance':
    pcs.explained_variance_ratio_, 'Cumulative proportion':
    np.cumsum(pcs.explained_variance_ratio_)})
pcsSummary_df = pcsSummary_df.transpose()
pcsSummary_df.columns = ['PC{}'.format(i) for i in range(1,
    len(pcsSummary_df.columns) + 1)]
pcsSummary_df.round(4)
```

	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	83.7641	70.9143	22.6437	19.1815	8.4232	2.0917
Proportion of variance	0.5395	0.3867	0.0394	0.0283	0.0055	0.0003
Cumulative proportion	0.5395	0.9262	0.9656	0.9939	0.9993	0.9997

	PC7	PC8	PC9	PC10	PC11	PC12	PC13
Standard deviation	1.6994	0.7796	0.6578	0.3704	0.1864	0.063	0.0
Proportion of variance	0.0002	0.0000	0.0000	0.0000	0.0000	0.000	0.0
Cumulative proportion	0.9999	1.0000	1.0000	1.0000	1.0000	1.000	1.0



The first two components account for 93% of the total variance, so using 2-3 components in further modeling would probably be sufficient

The Weightings for the First Five Components

```
pd.DataFrame(pcs.components_, columns=pcsSummary_df.columns,  
             index=cereal_df.iloc[:,3:].columns)  
pcsComponents_df.iloc[:, :5]
```

	PC1	PC2	PC3	PC4	PC5
calories	-0.077984	-0.009312	0.629206	-0.601021	0.454959
protein	0.000757	0.008801	0.001026	0.003200	0.056176
fat	0.000102	0.002699	0.016196	-0.025262	-0.016098
sodium	-0.980215	0.140896	-0.135902	-0.000968	0.013948
fiber	0.005413	0.030681	-0.018191	0.020472	0.013605
carbo	-0.017246	-0.016783	0.017370	0.025948	0.349267
sugars	-0.002989	-0.000253	0.097705	-0.115481	-0.299066
potass	0.134900	0.986562	0.036782	-0.042176	-0.047151
vitamins	-0.094293	0.016729	0.691978	0.714118	-0.037009
shelf	0.001541	0.004360	0.012489	0.005647	-0.007876
weight	-0.000512	0.000999	0.003806	-0.002546	0.003022
cups	-0.000510	-0.001591	0.000694	0.000985	0.002148
rating	0.075296	0.071742	-0.307947	0.334534	0.757708

Generalization

$X_1, X_2, X_3, \dots, X_p$, original p variables

$Z_1, Z_2, Z_3, \dots, Z_p$, weighted averages of original variables

All pairs of Z variables have 0 correlation

Order Z 's by variance (Z_1 largest, Z_p smallest)

Usually the first few Z variables contain most of the information, and so the rest can be dropped.

Normalizing data

- In these results, sodium dominates first PC
- Just because of the way it is measured (mg), its scale is greater than almost all other variables
- Hence its variance will be a dominant component of the total variance
- Normalize each variable to remove scale effect
Divide by std. deviation (may subtract mean first)
- Normalization (= standardization) is usually performed in PCA; otherwise measurement units affect results

```
from sklearn.preprocessing import scale  
pcs = PCA()  
pcs.fit(preprocessing.scale(cereals_df.iloc[:, 3:].dropna()))
```



Normalize the variables

PCA Output Using all 13 *Normalized* Numerical Variables

	PC1	PC2	PC3	PC4	PC5	PC6	
Standard deviation	1.9192	1.7864	1.3912	1.0166	1.0015	0.8555	
Proportion of variance	0.2795	0.2422	0.1469	0.0784	0.0761	0.0555	
Cumulative proportion	0.2795	0.5217	0.6685	0.7470	0.8231	0.8786	
	PC7	PC8	PC9	PC10	PC11	PC12	PC13
Standard deviation	0.8251	0.6496	0.5658	0.3051	0.2537	0.1399	0.0
Proportion of variance	0.0517	0.0320	0.0243	0.0071	0.0049	0.0015	0.0
Cumulative proportion	0.9303	0.9623	0.9866	0.9936	0.9985	1.0000	1.0

Weightings for the First Five *Normalized* Components

	PC1	PC2	PC3	PC4	PC5
calories	-0.299542	-0.393148	0.114857	-0.204359	0.203899
protein	0.307356	-0.165323	0.277282	-0.300743	0.319749
fat	-0.039915	-0.345724	-0.204890	-0.186833	0.586893
sodium	-0.183397	-0.137221	0.389431	-0.120337	-0.338364
fiber	0.453490	-0.179812	0.069766	-0.039174	-0.255119
carbo	-0.192449	0.149448	0.562452	-0.087835	0.182743
sugars	-0.228068	-0.351434	-0.355405	0.022707	-0.314872
potass	0.401964	-0.300544	0.067620	-0.090878	-0.148360
vitamins	-0.115980	-0.172909	0.387859	0.604111	-0.049287
shelf	0.171263	-0.265050	-0.001531	0.638879	0.329101
weight	-0.050299	-0.450309	0.247138	-0.153429	-0.221283
cups	-0.294636	0.212248	0.140000	-0.047489	0.120816
rating	0.438378	0.251539	0.181842	-0.038316	0.057584

- To convert the initial variables of a new observation to the PCs, you first normalize the new observation using the same normalization method applied to the training data. Then, you apply the PCA transformation learned from the training set (using the eigenvectors of the covariance matrix) to the normalized new data to obtain the PCs.
- [Video](#)

PCA in Classification/Prediction

- Apply PCA to training data
- Decide how many PC's to use
- Use variable weights in those PC's with validation/new data
- This creates a new reduced set of predictors in validation/new data

Regression-Based Dimension Reduction

- Multiple Linear Regression or Logistic Regression
- Use subset selection
- Algorithm chooses a subset of variables
- This procedure is integrated directly into the predictive task